PROJECT SUMMARY · LOOMI CONNECT AI HACKATHON

Scope Scout: An Agent That Reasons Like a Senior Implementation Architect

The Scope Scouts · Track 2: Engagement Intelligence & Workspace Diagnostics · 2026-05-28

The problem. Scoping a Bloomreach Engagement implementation is the most expertise-bottlenecked moment in a partner engagement. When a client hands an architect a list of desired triggers (twenty, thirty, sometimes more), the architect has to hold three knowledge layers in their head at once: the Bloomreach Plug & Play use case library and what's productized in it; the data model requirements for each use case (which attributes and events must exist); and pattern recognition across the request set itself, distinguishing structural duplicates, integration-dependent items, and vague client language from clean library matches. Senior architects do this work fluently. Junior architects can't yet. The result is inconsistent scopes, missed dependencies, and tacit expertise locked in a few heads.

Who Scope Scout is for. Solution architects at Bloomreach SI partners, at the moment of receiving a client's stated trigger list and producing the discovery agenda that will drive the next client conversation.

What Scope Scout does. Scope Scout reads a client's trigger list and reasons over it the way a senior architect does: forming provisional classifications against the published Plug & Play library, recognizing patterns across the request set (structural duplicates, shared dependencies, lifecycle gaps), and distinguishing language the library can resolve from language only the client can. The agent holds its reads provisionally and revises them as workspace evidence and contextual cues accumulate, propagating consequences through related classifications. The output is a structured discovery agenda with a fully visible reasoning trace; every classification, dependency flag, and recommendation includes the evidence and inference behind it. Architects review; Sawyer proposes.

How it works. Scope Scout runs an agentic reasoning loop on Anthropic's `claude-sonnet-4-6`: the agent iteratively calls live workspace-introspection tools (Loomi Connect Marketing MCP) and local Plug & Play knowledge tools, reasoning over each result until the discovery agenda is complete, bounded at 25 tool iterations for safety. The two knowledge sources are Bloomreach's published use case library and the live workspace state retrieved through the Marketing MCP. The reasoning is adaptive: later evidence revises earlier classifications, and consequences propagate across related findings.

Why this matters. Scoping is the first billable touchpoint with a new client. Errors compound: they undermine trust, delay time-to-value, and create downstream delivery risk. By externalizing the tacit pattern recognition that senior architects do without thinking, Scope Scout lets junior architects produce senior-architect-quality scopes consistently, accelerates the partner channel's throughput of clean implementations, and demonstrates how Loomi Connect can serve as agentic infrastructure for partner enablement, not only direct customer use.

What it isn't. Scope Scout is a reasoning partner, not a pricing tool. It produces the discovery agenda that becomes input to pricing; hour math stays a human-driven step downstream. All outputs are recommendations for architect review before client delivery. Inspection is read-only. The orienting fields in the console UI help structure the architect's input; the agent reasons over the combined brief. Demo uses fixture data; the reasoning logic is identical to production.